

CompSys Review



Understanding the OS, Races, VM Address Translation

TA: Julian Philippe Pedersen <jupe@di.ku.dk>

CompSys 2022-2023

Understanding the OS

Kernel

- Kernel is a protected part of the OS.
- Only the kernel has direct access to:
 - Hardware
 - System memory
 - File I/O
- System calls require access to the kernel.

PROGRAMS VS. PROCESSES

- A program is just a list of instructions for the computer. (A file containing code)
- A process is an instance of a program. This instance is “live” and stateful.
 - When people say “program,” they often mean “process.”
- When multiple processes are running, the OS utilizes context switching to give the illusion of concurrency.
 - Context switching requires the saving of the state of a process. To do this, it must:
 - All registers (including control registers)
 - Contents of memory

Races

- Race conditions: a situation where values depend on order of execution of instructions
 - Problem in concurrent execution
 - Multithreaded solutions: semaphores (mutex), conditional variables
 - Less of a problem with threads (but race conditions can still exist!)

Threads

- Run in the same address space as the calling process
- Threads have their own thread context, which includes:
 - Thread ID
 - Stack + Stack pointer
 - Program counter
 - General-purpose registers
 - Condition codes

Dealing with Thread Concurrency

- Semaphores (mutexes)
 - Invariant: must always be non-negative
 - `P()` // Prolaag, "try". Decrements semaphore, unless decrement would cause semaphore to become < 0 .
 - `V()` // Verhoog, "increment". Increments semaphore.
- Conditional variables can allow blocking and signaling specific threads
 - Watch out for spurious wakeups! (When a thread continues execution even when the valid condition is not true.)
 - When in doubt, while() it out!
 - This ensures that when a thread is signaled to wake up, it will always check if the condition for it to continue is true.
 - If-statements cause this check to only occur once during a thread's lifetime. This is not necessarily incorrect, but it can be difficult to guard against spurious wakeups.

Processes

- Processes contain a duplicate of the parent's information, but in a separate address space.
 - Changes made in one process are not reflected in the others!
- Child processes must be reaped by their parent.
 - Processes not reaped by the parent upon parent termination will be “adopted” by the `init` process (pid 1).
 - Long-running child processes adopted by `init` that never terminate are called “zombie” processes.

Fork ()

- `Fork ()` returns two values
 - Child's process id (PID) in parent process
 - 0 in child process
 - Use `fork ()` return value to check whether the current program is a parent or child process
 - `fork () == 0`
 - `True`: We are in the parent process
 - `False`: We are in the child process
- `waitpid(pid, *status, options)` is used to block execution in parent processes based on child processes
 - `(waitpid(child_pid, status, 0) > 0)` blocks until child with `child_pid` is finished executing

Fork() cont.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <sys/types.h>
5  #include <sys/wait.h>
6
7  int main(void) {
8      pid_t some_pid;
9      if((some_pid = fork()) == 0) {          // Kør fork(), og gem fork() output i some_pid
10         printf("I'm a child! fork() returned: %d\n", some_pid); // Denne code kører kun i child process, og some_pid vil være 0 her
11         exit(0);
12     } else {
13         while (waitpid(some_pid, NULL, 0) > 0); // Vent på, at barnet kører færdigt
14         printf("I'm a parent! fork() returned: %d\n", some_pid); // Denne code kører kun i parent process, og some_pid vil være barnets pid her
15     }
16     return EXIT_SUCCESS;
17 }
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

```
● julianpedersen@DESKTOP-6200G9A:/mnt/c/Users/J/Documents/KU/test$ ./forkexample
I'm a child! fork() returned: 0
I'm a parent! fork() returned: 11777
○ julianpedersen@DESKTOP-6200G9A:/mnt/c/Users/J/Documents/KU/test$
```

Process Graphs

- Very helpful in determining execution order of forked processes!
- Things to keep in mind:
 - `fork()` spawns two edges
 - `wait()` can “join” multiple edges into a single one (watch for child execution order!)
 - `exit()` causes the calling process to exit immediately. It will not continue execution of commands below.

Fork() Example

```
int main () {  
    if (fork() == 0) {  
        printf("1");  
        if (fork() == 0) {  
            printf("2");  
        } else {  
            pid_t pid; int status;  
            if ((pid = waitpid(pid, NULL, 0)) > 0) {  
                printf("3");  
            }  
        }  
    } else {  
        printf("4");  
        exit(0);  
    }  
    printf("5");  
}
```


```
int main () {  
    if (fork() == 0) {  
        printf("1");  
        if (fork() == 0) {  
            printf("2");  
        } else {  
            pid_t pid; int status;  
            if ((pid = waitpid(pid, NULL, 0)) > 0) {  
                printf("3");  
            }  
        }  
    } else {  
        printf("4");  
        exit(0);  
    }  
    printf("5");  
}
```

Handwritten annotations:

- child* (next to the first `if (fork() == 0)`)
- child-child* (next to the second `if (fork() == 0)`)
- wait until child-child has finished executing!* (next to the `waitpid` call)
- parent* (next to the `else` block)
- Exit() stops parent here!* (next to the `exit(0)` call)
- Only child and child-child reach this line!* (next to the final `printf("5");`)

Fork() Process Graph

```
int main () {  
    if (fork() == 0) {  
        printf("1");  
        if (fork() == 0) {  
            printf("2");  
        } else {  
            pid_t pid; int status;  
            if ((pid = waitpid(pid, NULL, 0)) > 0) {  
                printf("3");  
            }  
        }  
    } else {  
        printf("4");  
        exit(0);  
    }  
    printf("5");  
}
```



VM Address Translation

Background

- VPA (Virtual Page Address)
 - VPA can be split into Virtual Page Number (VPN) and Virtual Page Offset (VPO)
 - VPN can be further split into Virtual Page Tag (VPT) and Virtual Page Index (VPI)
- PPA (Physical Page Address)
 - PPA can be split into Physical Page Number (PPN) and Physical Page Offset (PPO)
- TLB (Translation Lookaside Buffer)
 - Contains previously-accessed VPAs/PPAs by storing VPTs, PPNs, valid bits in “ways”
 - Analogous to cache
- Page Table
 - Contains Page Table Entries (PTEs)
 - PTEs consist of VPNs, PPNs, valid bits

Tips

- Read the description!
 - ☆ Look for keywords **page size**, **# sets**, **address format**
 - # lower-order bits = $\log_2(\text{page size})$
 - # set index bits = $\log_2(\text{\# sets})$
 - Address format is almost always hexadecimal, but make sure to check! (do not confuse yourself between hexadecimal and base-10!)
- When you have found the number of bits for VPT/VPI/VPO/ and PPN/PPO, mark the ranges down on the test!
 - Easy to mess this up, so the less data in your “working set” (short-term memory) the better
- Each hexadecimal digit corresponds to 4 binary bits in the same order.
 - $0x00 = 0000\ 0000$, $0x01 = 0000\ 0001$... $0x10 = 0001\ 0000$...
 $0xff = 1111\ 1111$
 - Useful for quick translation between binary and hexadecimal

"Algorithm"

- Read description. Find page size and # of sets.
 - Determine number of bits for (V/P)PO with $\log_2(\text{page size})$ and TLBI with $\log_2(\text{\# sets})$.
- Translate the given address from hex to binary and write the bits in the top bitfield.
- Translate VPN, TLB index (TLBI), TLB tag (TLBT).
- Lookup TLB with TLBI and TLBT
 - If (tag exists in index) && (valid bit = 1), TLB hit 😊
 - Copy VPO to PPO directly. Copy PPN from TLB to PPN field, right-to-left. You're done! ✓
 - If (tag does not exist in index) || (valid bit = 0), TLB miss ☹️ Continue to next step.
- Lookup page table with VPN
 - If (VPN exists in page table) && (valid bit = 1), no page fault! 😊
 - Copy VPO to PPO directly. Copy PPN from PTE to PPN field, right-to-left. You're done! ✓
 - If (VPN exists in page table) || (valid bit = 0), page fault ☹️ Continue to next step
- If (no TLB hit) && (page fault)
 - No valid PPN exists. Leave bottom bitfield empty. You're done! ✓

Example: TLB Hit

Bit position	13	12	11	10	9	8	7	6	5	4	3	2	1	0
VA = 0x03d4	0	0	0	0	1	1	1	1	0	1	0	1	0	0
	0x0				0x3				0xd				0x4	

Each hex cipher corresponds to 4 bits in binary.

	TLBT								TLBI						
	<u>0x03</u>								<u>0x3</u>						
	0x0		0x3						0x3						
Bit position	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
VA = 0x03d4	0	0	0	0	1	1	1	1	0	1	0	1	0	0	
	0x0					0xf			0x1		0x4				
	0x0f								<u>0x14</u>						
	VPN								VPO						

1. Extracting VPN, TLBT, TLBI, VPO from a VA.

	0xA				0x6				0x4			
Bit position	11	10	9	8	7	6	5	4	3	2	1	0
PA = 0xA64	1	0	1	0	0	1	0	1	0	1	0	0
	0x2		0x9				0x1		0x4			
	<u>0x29</u>							<u>0x14</u>				
	PPN							PPO				

3. Constructing PA from PPA and PPO.

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	0x03	—	0	0x09	0x0D	1	0x00	—	0	0x07	0x02	1
1	0x03	0x2D	1	0x02	—	0	0x04	—	0	0x0A	—	0
2	0x02	—	0	0x08	—	0	0x06	—	0	0x03	—	0
3	0x07	—	0	<u>0x03</u>	<u>0x29</u>	1	0x0A	0x34	1	0x02	—	0

2. Looking up 0x03 in a TLB.